

# The Art of Incremental Stream Processing

@markhibberd

I have a problem

q:Quit   d:Del   u:Undel   s:Save   m:Mail				
INBOX [52]	10/6422		[ +]	04/18
archive	34982/52900		[ +]	04/18
sent [18]	38/ 840		[ +]	04/18
drafts	1/ 54		[ F]	04/18
			[ !]	04/18
			[ +]	04/18
			[ ]	04/18
			[ ]	04/18

**~150k Emails**

~4GB Emails

**100s** of new email / day

**~5 I want to Read**

~2 I want to Reply To



Messages delivered where  
and when I need them

Ability to **locate** important  
messages from days past

From oleg-at-okmij.org Thu Sep 18 23:51:59 2008  
To: haskell-cafe@haskell.org  
Subject: Lazy vs correct IO [Was: A round of golf]  
Message-ID: <20080919065159.616A5AF09@Adric.metnet.fnmoc.navy.mil>  
Date: Thu, 18 Sep 2008 23:51:59 -0700 (PDT)  
Status: OR

Lennart Augustsson wrote

```
> main = do
>   name:_ <- getArgs
>   file <- readFile name
>   print $ length $ lines file
```

Given the stance against top-level mutable variables, I have not expected to see this Lazy IO code. After all, what could be more against the spirit of Haskell than a `pure' function with observable side effects. With Lazy IO, one indeed has to choose between correctness and performance. The appearance of such code is especially strange after the evidence of deadlocks with Lazy IO, presented on this list less than a month ago. Let alone unpredictable resource usage and reliance on finalizers to close files (forgetting that GHC does not guarantee that finalizers will be run at all).

Is there an alternative?

From oleg-at-okmij.org Thu Sep 18 23:51:59 2008  
To: haskell-cafe@haskell.org  
Subject: Lazy vs correct IO [Was: A round of golf]  
Message-ID: <20080919065159.616A5AF09@Adric.metnet.fnmoc.navy.mil>  
Date: Thu, 18 Sep 2008 23:51:59 -0700 (PDT)  
Status: OR

Lennart Augustsson wrote

```
> main = do
>   name:_ <- getArgs
>   file <- readFile name
>   print $ length $ lines file
```

Given the stance against top-level mutable variables, I have not expected to see this Lazy IO code. After all, what could be more against the spirit of Haskell than a `pure' function with observable side effects. With Lazy IO, one indeed has to choose between correctness and performance. The appearance of such code is especially strange after the evidence of deadlocks with Lazy IO, presented on this list less than a month ago. Let alone unpredictable resource usage and reliance on finalizers to close files (forgetting that GHC does not guarantee that finalizers will be run at all).

Is there an alternative?

Zeitgeist

```
13 find "$MAILDIR" -type f           | \
14   xargs -n 1 stat -f "%m|%N"      | \
15   sort -n                          | \
16   cut -d '|' -f 2                  | \
17   xargs grep -l "$QUERY"           | \
18   head -5                          | \
19   xargs less
```

```
5 data Email =  
6   Email { date :: Int,  content :: String }  
7   deriving (Show, Eq)  
8  
9 search :: String -> [Email] -> [Email]  
10 search term =  
11   take 5  
12   . filter (isInfixOf term . content)  
13   . sortBy (compare `on` date)
```

**Reality** Calling



“--mmap

Use mmap(2) instead of read(2) to read input, which can result in better performance under some circumstances but can cause undefined behaviour.”

— \$(man grep)

```

275 struct file *
276 grep_open(const char *path)
277 {
278     struct file *f;
279
280     f = grep_malloc(sizeof *f);
281     memset(f, 0, sizeof *f);
282     if (path == NULL) {
283         /* Processing stdin implies --line-buffered.
284         */
285         lbflag = true;
286         f->fd = STDIN_FILENO;
287     } else if ((f->fd = open(path, 0_RDONLY)) == -
288               1)
289         goto error1;
290
291     if (filebehave == FILE_MMAP) {
292         struct stat st;
293
294         if ((fstat(f->fd, &st) == -1) || (st.st_size >
295     OFF_MAX) ||
296         (!S_ISREG(st.st_mode)))
297             filebehave = FILE_STDIO;
298         else {
299             int flags = MAP_PRIVATE | MAP_NOCORE |
300                 MAP_NOSYNC;
301 #ifdef MAP_PREFAULT_READ
302             flags |= MAP_PREFAULT_READ;
303 #endif
304             fsiz = st.st_size;
305             buffer = mmap(NULL, fsiz, PROT_READ, flags,
306                 f->fd, (off_t)0);
307             if (buffer == MAP_FAILED)
308                 filebehave = FILE_STDIO;
309             else {
310                 bufrem = st.st_size;
311                 bufpos = buffer;
312                 madvise(buffer, st.st_size, MADV_SEQUENTIAL)
313             ;
314             }
315         }
316     }
317
318     if ((buffer == NULL) || (buffer == MAP_FAILED))
319         buffer = grep_malloc(MAXBUFFSIZ);
320
321     if (filebehave == FILE_GZIP &&
322         (gzbufdesc = gzdopen(f->fd, "r")) == NULL)
323         goto error2;
324
325 #ifndef WITHOUT_BZIP2
326     if (filebehave == FILE_BZIP &&
327         (bzbufdesc = BZ2_bzdopen(f->fd, "r")) ==
328         NULL)
329         goto error2;
330 #endif
331
332     /* Fill read buffer, also catches errors early
333     */
334     if (bufrem == 0 && grep_refill(f) != 0)
335         goto error2;
336
337     /* Check for binary stuff, if necessary */
338     if (binbehave != BINFILE_TEXT && memchr(bufpos,
339         '\0', bufrem) != NULL)
340         f->binary = true;
341
342     return (f);
343
344 error2:
345     close(f->fd);
346 error1:
347     free(f);
348     return (NULL);
349 }
350
351

```

**“With Lazy IO, one indeed has  
to choose between correctness  
and performance.”**

**— Oleg Kiselyov**

```
5 type Maildir =  
6   FilePath  
7  
8 data Email =  
9   Email { date :: Int,   content :: String }  
10  deriving (Show, Eq)  
11  
12 search :: String -> Maildir -> IO [Email]  
13 search term =  
14   {- oh noes! It's so horrible  
15      I can't even show it -}
```

Is there an alternative?

# Intuition 1: A Language

I Need To Produce Values

```
1 type In i
```



I Need To Consume Values

```
1 type In i
2
3 type Out o
```

I Need To Transform Values

```
1 type In i
2
3 type Out o
4
5 data Pipeline i o
```

I *May* Have **Effects**

```
1 type In i m
2
3 type Out o m
4
5 data Pipeline i o m
```

I *May* **Compute** A Value

```
1 type In i m a
2
3 type Out o m a
4
5 data Pipeline i o m a
```



# A (Simple) Interface

```
1 type In i m a = Pipeline i () m a
2
3 type Out o m a = Pipeline Void o m a
4
5 data Pipeline i o m a
```

```
1 type In i m a = Pipeline i () m a
2
3 type Out o m a = Pipeline Void o m a
4
5 data Pipeline i o m a
6   = Done a
7   | Yield o (Pipeline i o a)
8   | Await (i -> Pipeline i o a)
```

# Intuition 2: Pipelines





Poll


Prepare

Classify

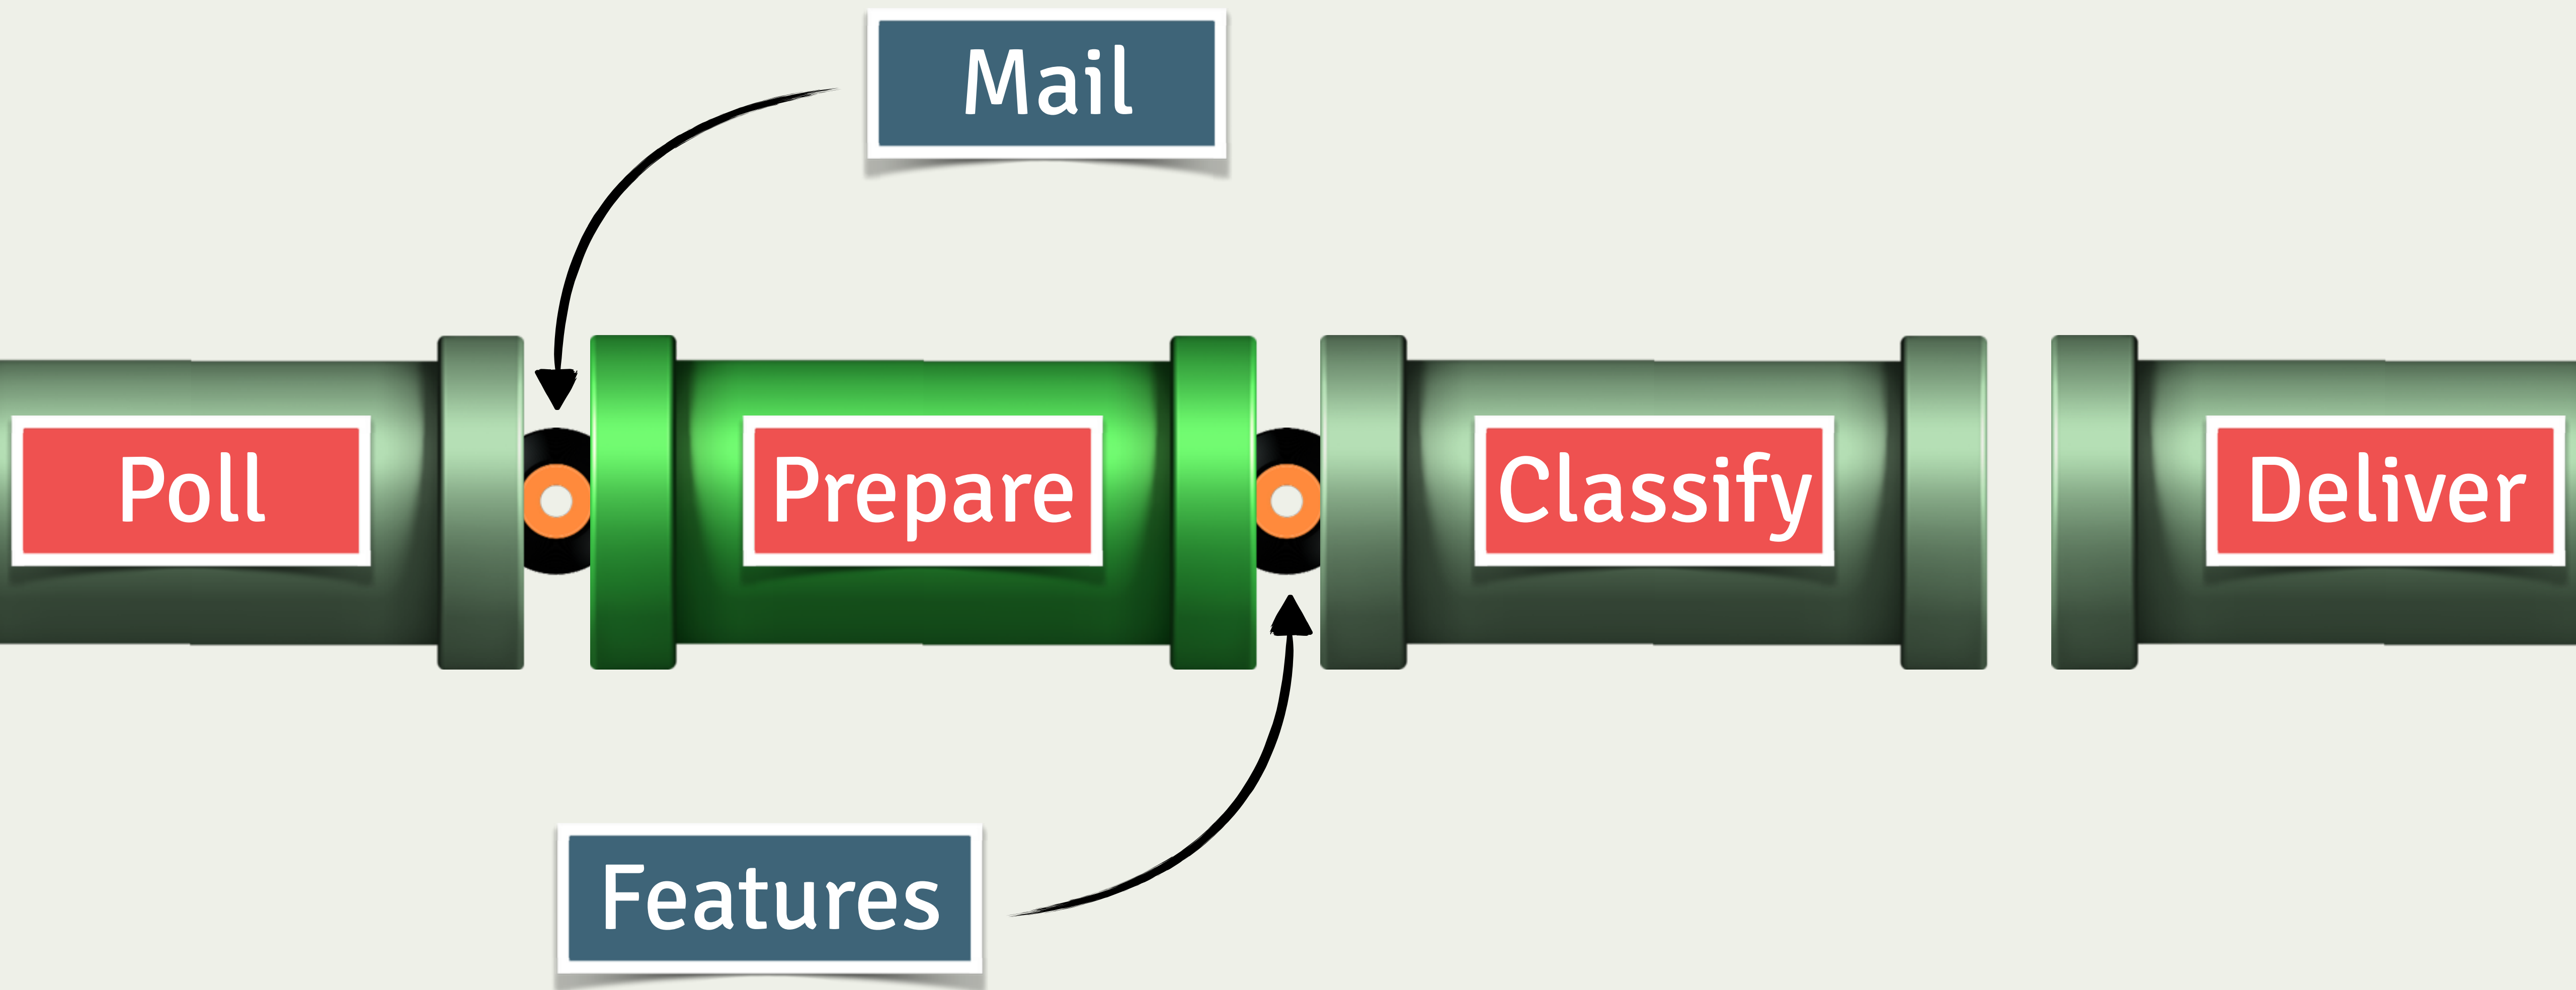
Deliver






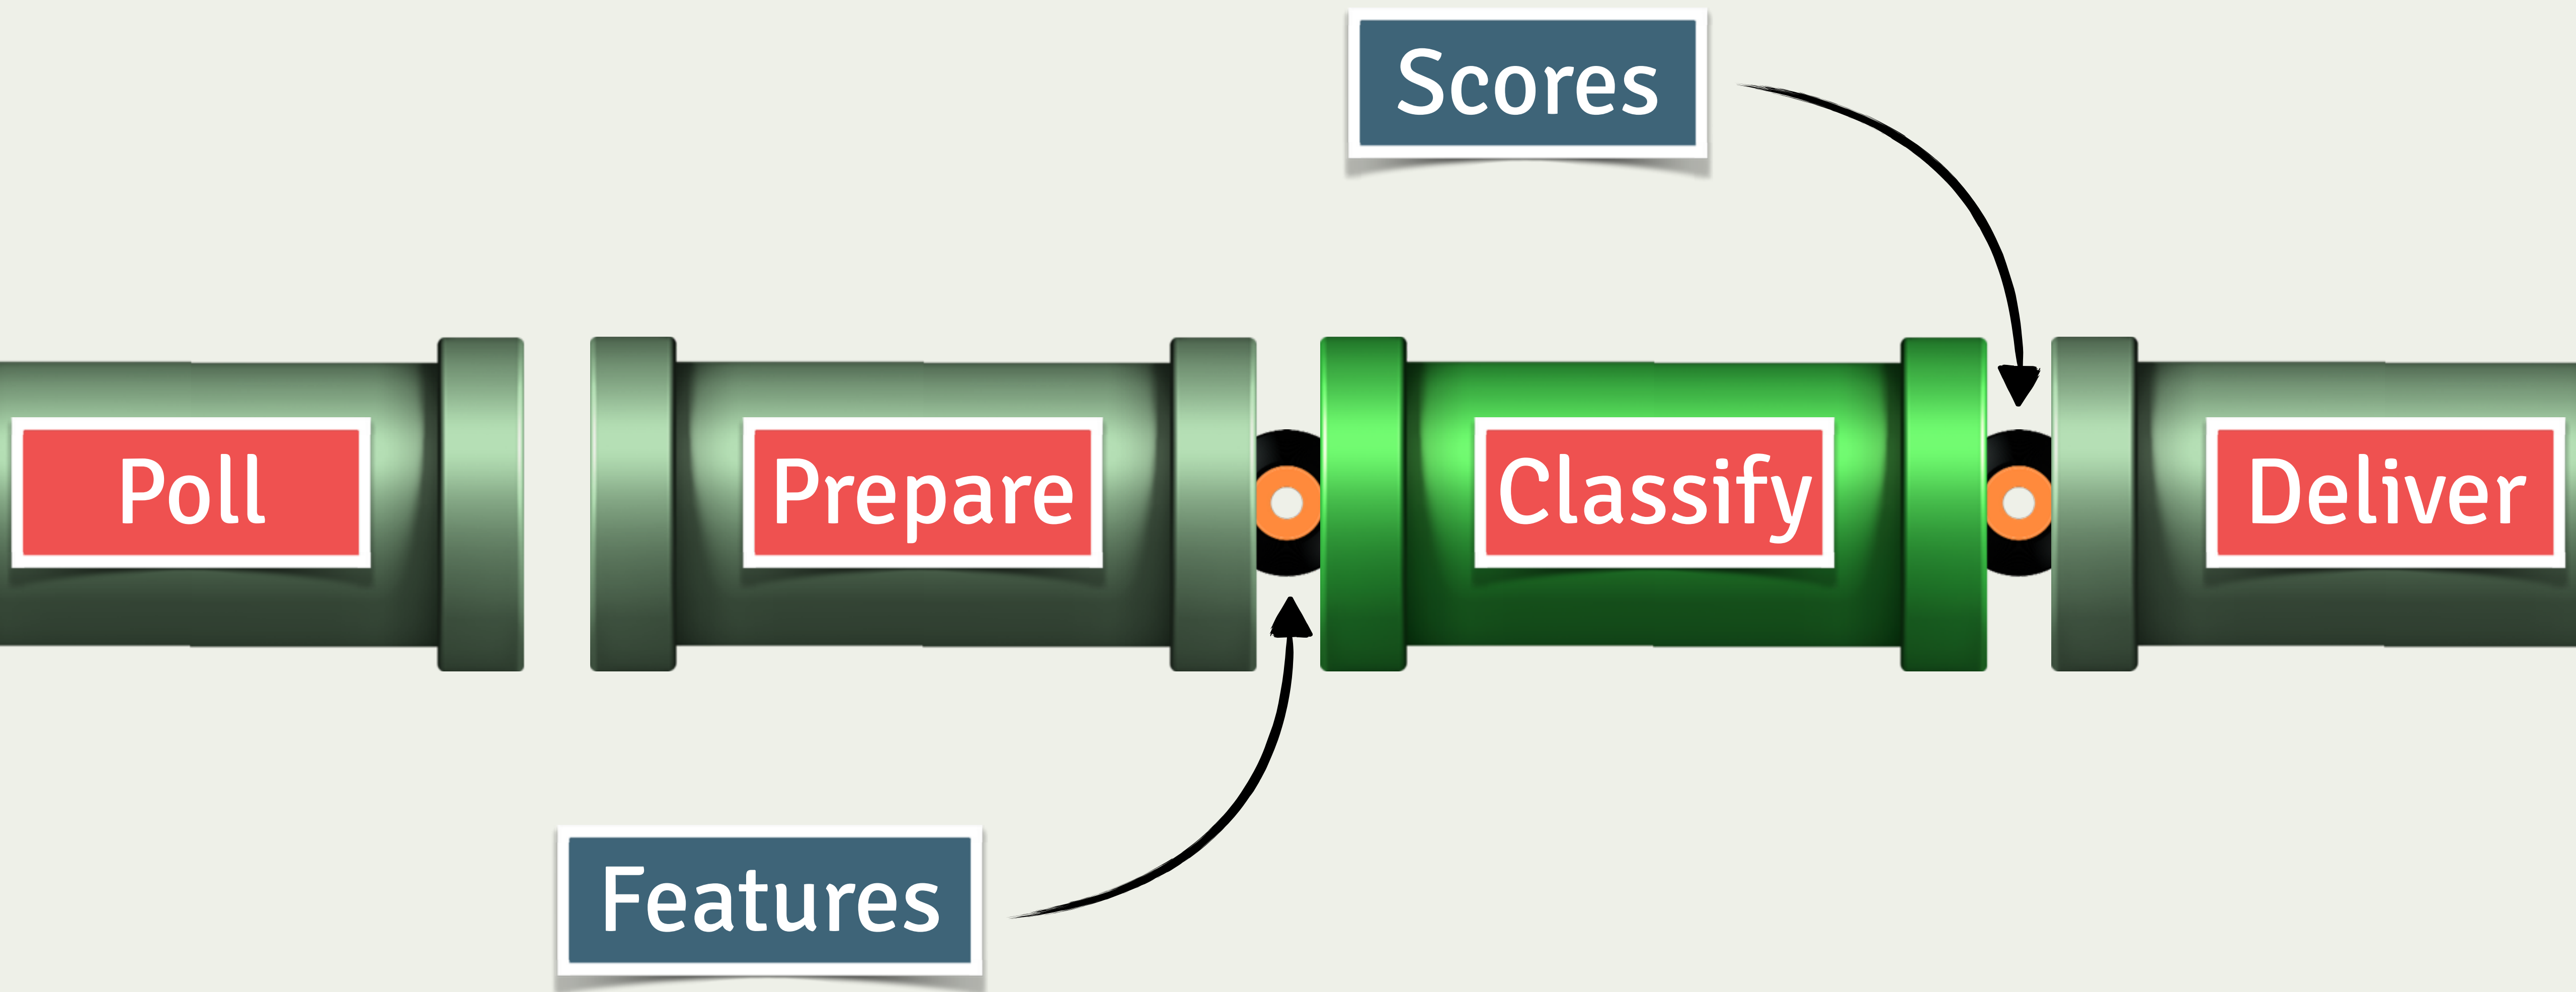


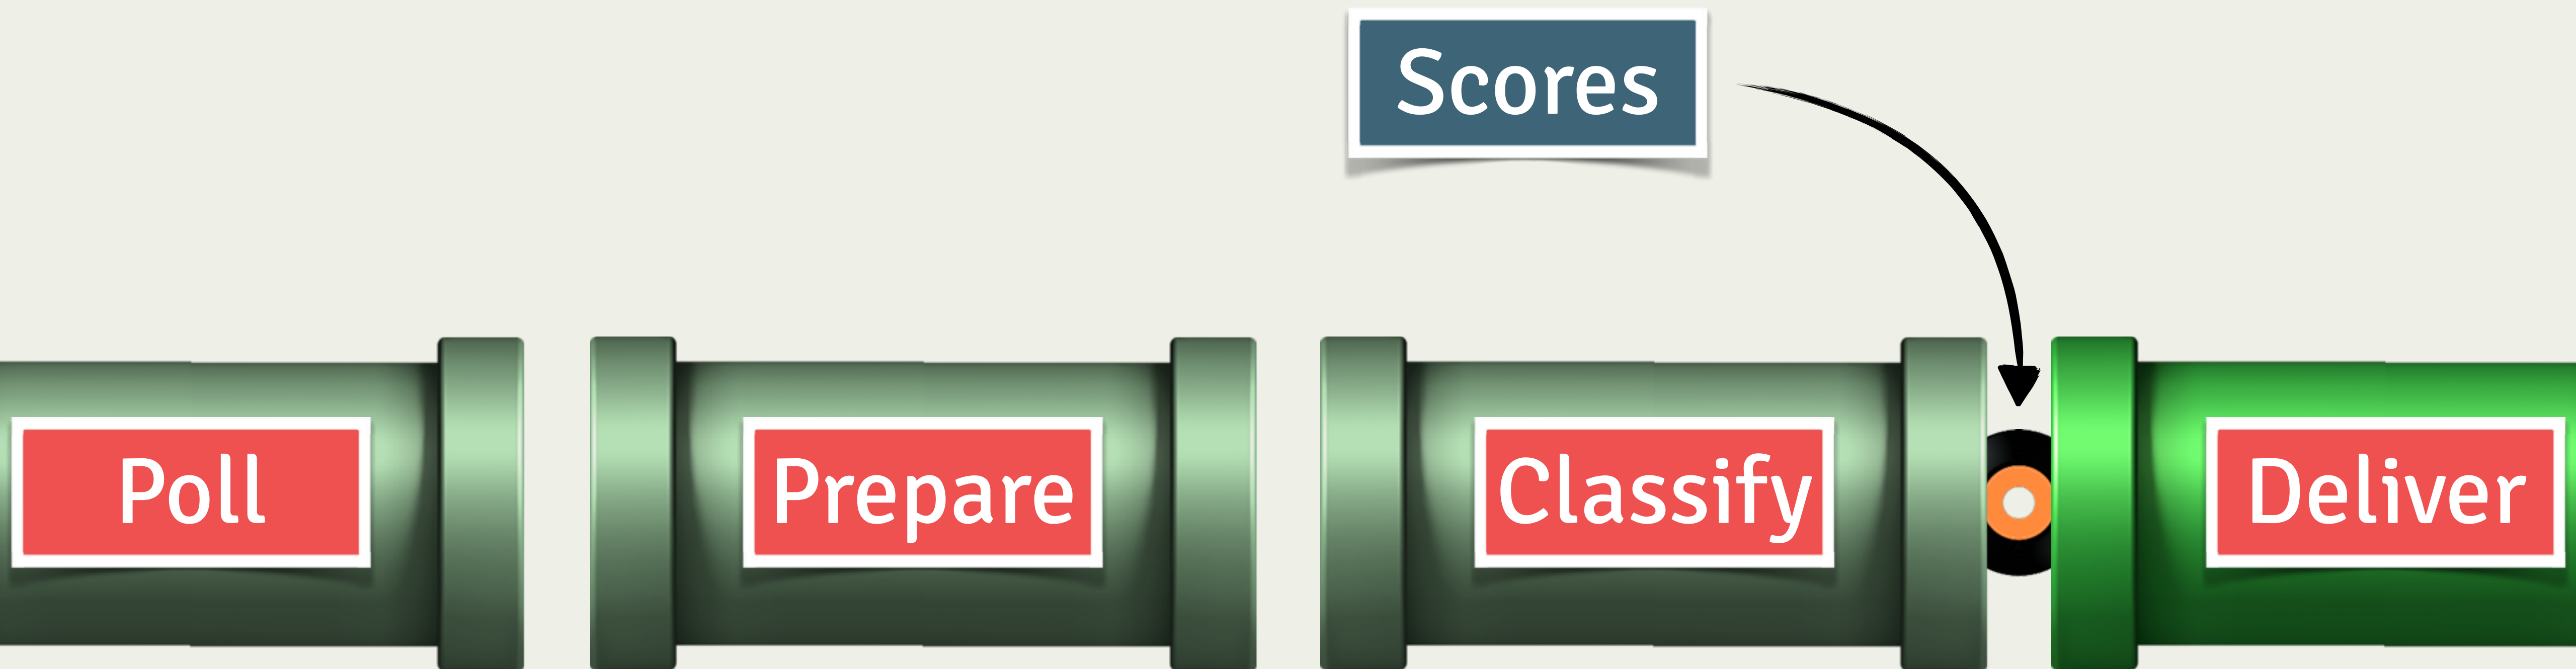







Features







Getting Real

```
1 data Proxy i' i o' o m a
2
3 type Producer i m a = Proxy X () () o m a
4 type Consumer o m a = Proxy () i () X m a
5 type Pipe i o m a    = Proxy () i () o m a
```

```
1 type In i m a
2
3 type Out o m a
4
5 type Pipeline i o m a
```

Pipes

```
1 data Proxy i' i o' o m a
2
3 type Producer i m a = Proxy X () () o m a
4 type Consumer o m a = Proxy () i () X m a
5 type Pipe i o m a    = Proxy () i () o m a
```

Explicit Input and Output  
at Each Component

Pipes



```
1 data Proxy i' i o' o m a
2
3 type Producer i m a = Proxy X () () o m a
4 type Consumer o m a = Proxy () i () X m a
5 type Pipe i o m a    = Proxy () i () o m a
```

Effects On Producers,  
Consumers And Pipes

Pipes

```
1 data Proxy i' i o' o m a
2
3 type Producer i m a = Proxy X () () o m a
4 type Consumer o m a = Proxy () i () X m a
5 type Pipe i o m a    = Proxy () i () o m a
```

Can Terminate With A  
Value Anywhere In Pipeline

Pipes

```
1 data Pipe l i o u m r
2
3 newtype ConduitM i o m r =
4   ConduitM { unConduitM :: Pipe i i o () m r }
5
6 type Source m o      = ConduitM () o      m ()
7 type Sink i m a      = ConduitM i  Void m a
8 type Conduit i m o    = ConduitM i  o      m ()
```

```
1 type In i m a
2
3 type Out o m a
4
5 type Pipeline i o m a
```

Conduit

```
1 data Pipe l i o u m r
2
3 newtype ConduitM i o m r =
4   ConduitM { unConduitM :: Pipe i i o () m r }
5
6 type Source m o      = ConduitM () o      m ()
7 type Sink i m a      = ConduitM i  Void m a
8 type Conduit i m o = ConduitM i  o      m ()
```

Explicit Input and Output  
at Each Component



Conduit

```
1 data Pipe l i o u m r
2
3 newtype ConduitM i o m r =
4   ConduitM { unConduitM :: Pipe i i o () m r }
5
6 type Source m o      = ConduitM () o      m ()
7 type Sink i m a      = ConduitM i  Void m a
8 type Conduit i m o    = ConduitM i  o      m ()
```

Effects On Sources,  
Sinks And Conduits



Conduit

```
1 data Pipe l i o u m r
2
3 newtype ConduitM i o m r =
4   ConduitM { unConduitM :: Pipe i i o () m r }
5
6 type Source m o      = ConduitM () o      m ()
7 type Sink i m a      = ConduitM i  Void m a
8 type Conduit i m o    = ConduitM i  o      m ()
```

Can Only Terminate With A  
Value On a Sink



Conduit

```
1 sealed abstract class Process[F[_],O]
2
3 type Process0[O]           = Process[Env[_,_]#Is, O]
4 type Process1[I, O]        = Process[Env[I,_]#Is, O]
5 type Sink[F[_], O]         = Process[F, O => F[Unit]]
```

```
1 type In i m a
2
3 type Out o m a
4
5 type Pipeline i o m a
```

Scalaz Streams

```
1 data Process m o
2
3 type Process0 o    = forall a. Process (Is a) o
4 type Process1 i o  = Process (Is i) o
5 type Sink m o      = Process m (o -> m ())
```

```
1 type In i m a
2
3 type Out o m a
4
5 type Pipeline i o m a
```

Scalaz Streams



```
1 data Process m o
2
3 type Process0 o    = forall a. Process (Is a) o
4 type Process1 i o  = Process (Is i) o
5 type Sink m o      = Process m (o -> m ())
```



Model Request And Production

Rather Than Input and Output

Scalaz Streams

```
1 data Process m o
2
3 type Process0 o    = forall a. Process (Is a) o
4 type Process1 i o  = Process (Is i) o
5 type Sink m o      = Process m (o -> m ())
```



Effects Are Returned As

Values, Transducers are Pure

Scalaz Streams


```
1 data Process m o
2
3 type Process0 o    = forall a. Process (Is a) o
4 type Process1 i o  = Process (Is i) o
5 type Sink m o      = Process m (o -> m ())
6
7 runFoldMap :: (Monad m, Monoid b) =>
8           Process m o -> (o -> m b) -> m b
```

Computation of Values  
Modelled Externally

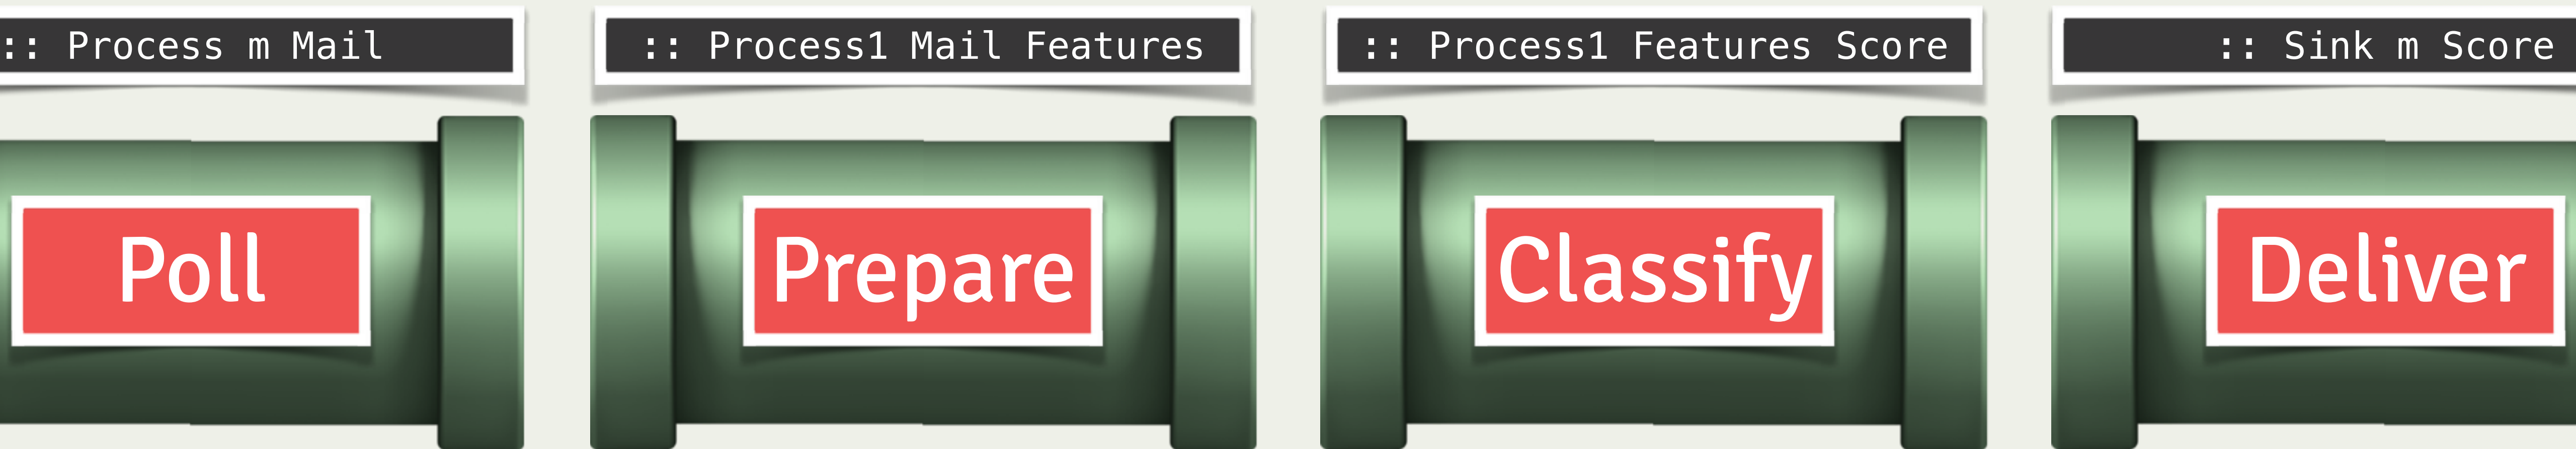
Scalaz Streams



Pipes



Conduit



Scalaz Stream

# Horizontal Composition

# Mail Delivery

Poll

Prepare

Classify

Deliver



Mail Delivery

Poll

Prepare

Classify


Deliver

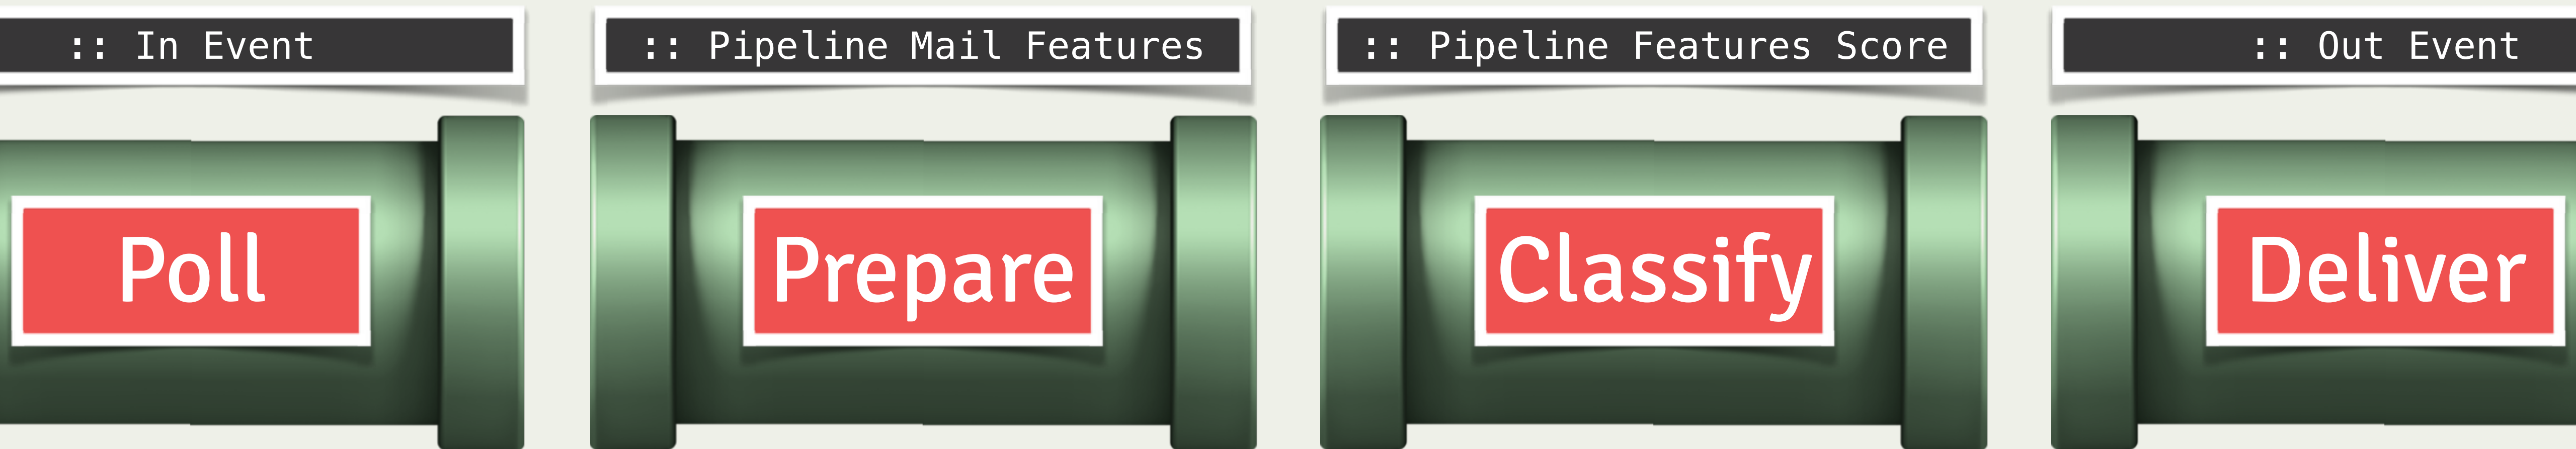
Poll

Scan

Throttle

Read





(>|) :: Pipeline i o m a -> Pipeline o o' m a -> Pipeline i o' m a

```
:: Pipeline () Void
```

```
:: In Event
```

```
:: Pipeline Mail Features
```

```
:: Pipeline Features Score
```

```
:: Out Event
```

Poll

>|

Prepare

>|

Classify

>|

Deliver

```
(>|) :: Pipeline i o m a -> Pipeline o o' m a -> Pipeline i o' m a
```

```
eval :: Pipeline () Void m a -> m a
```

```
:: In Event
```

```
:: Pipeline Mail Features
```

```
:: Pipeline Features Score
```

```
:: Out Event
```

Poll

>|

Prepare

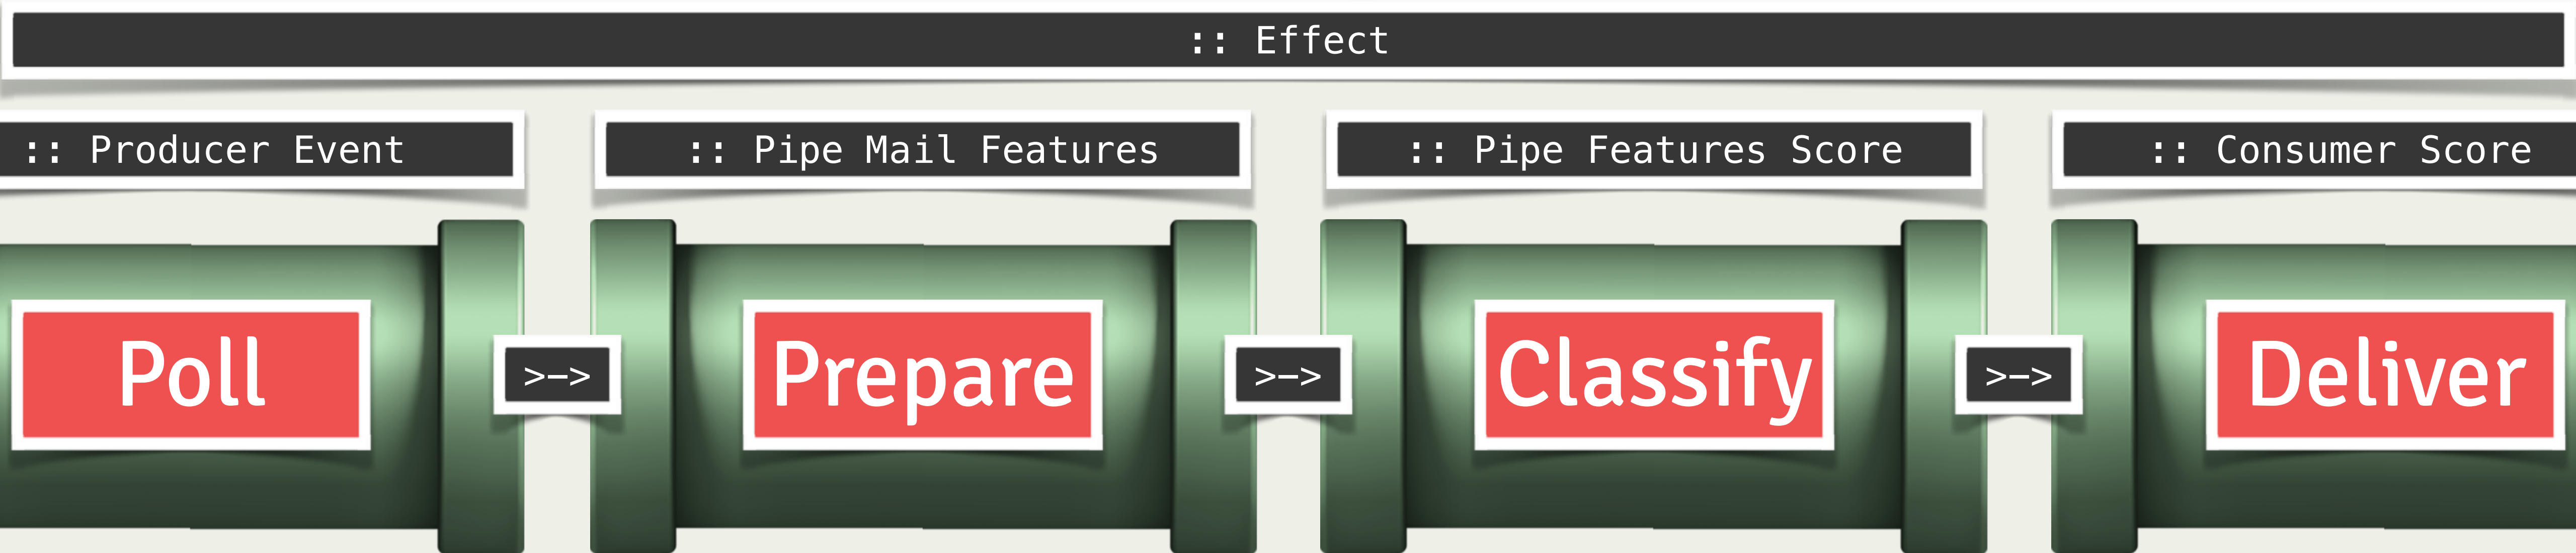
>|

Classify

>|

Deliver

```
(>|) :: Pipeline i o m a -> Pipeline o o' m a -> Pipeline i o' m a
```



Pipes



`runEffect :: Effect m a -> m a`

`:: Producer Event`

`:: Pipe Mail Features`

`:: Pipe Features Score`

`:: Consumer Score`

Poll

`>->`

Prepare

`>->`

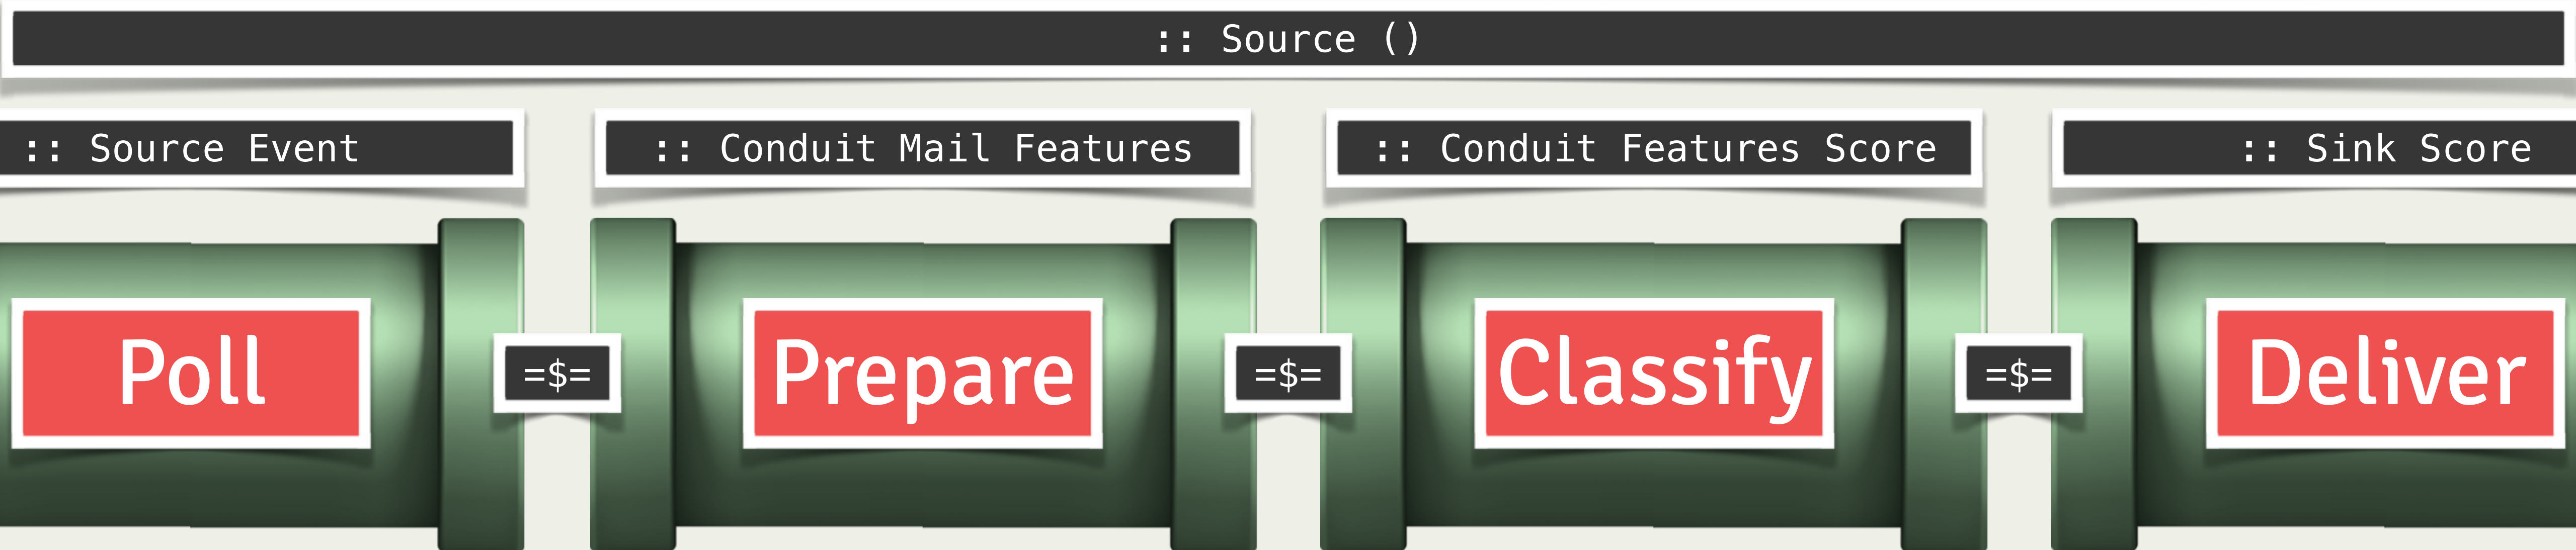
Classify

`>->`

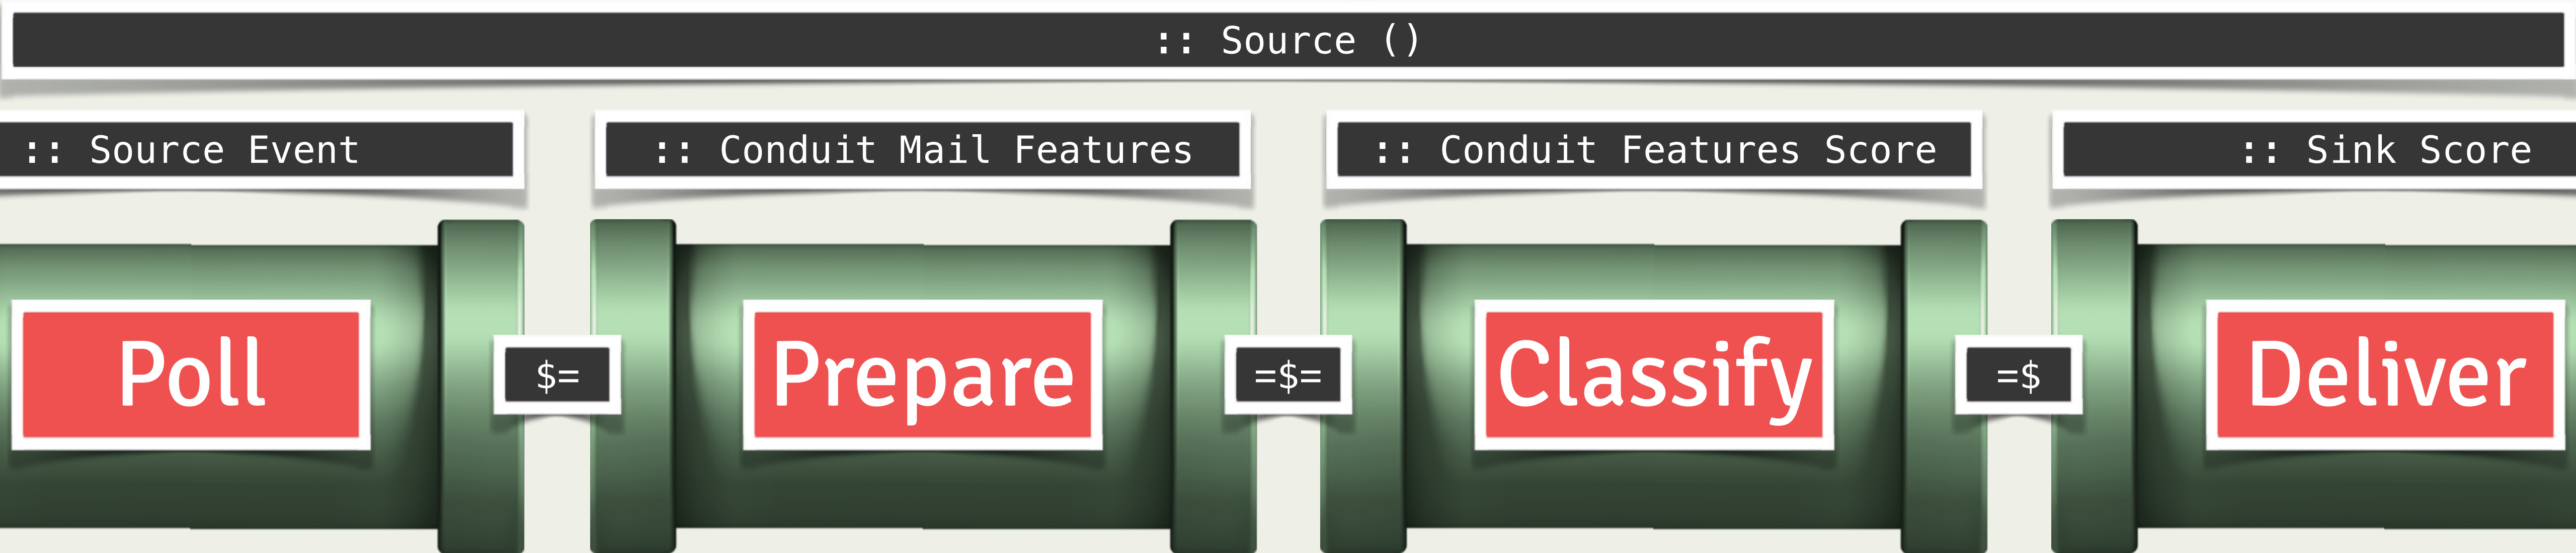
Deliver

Pipes'

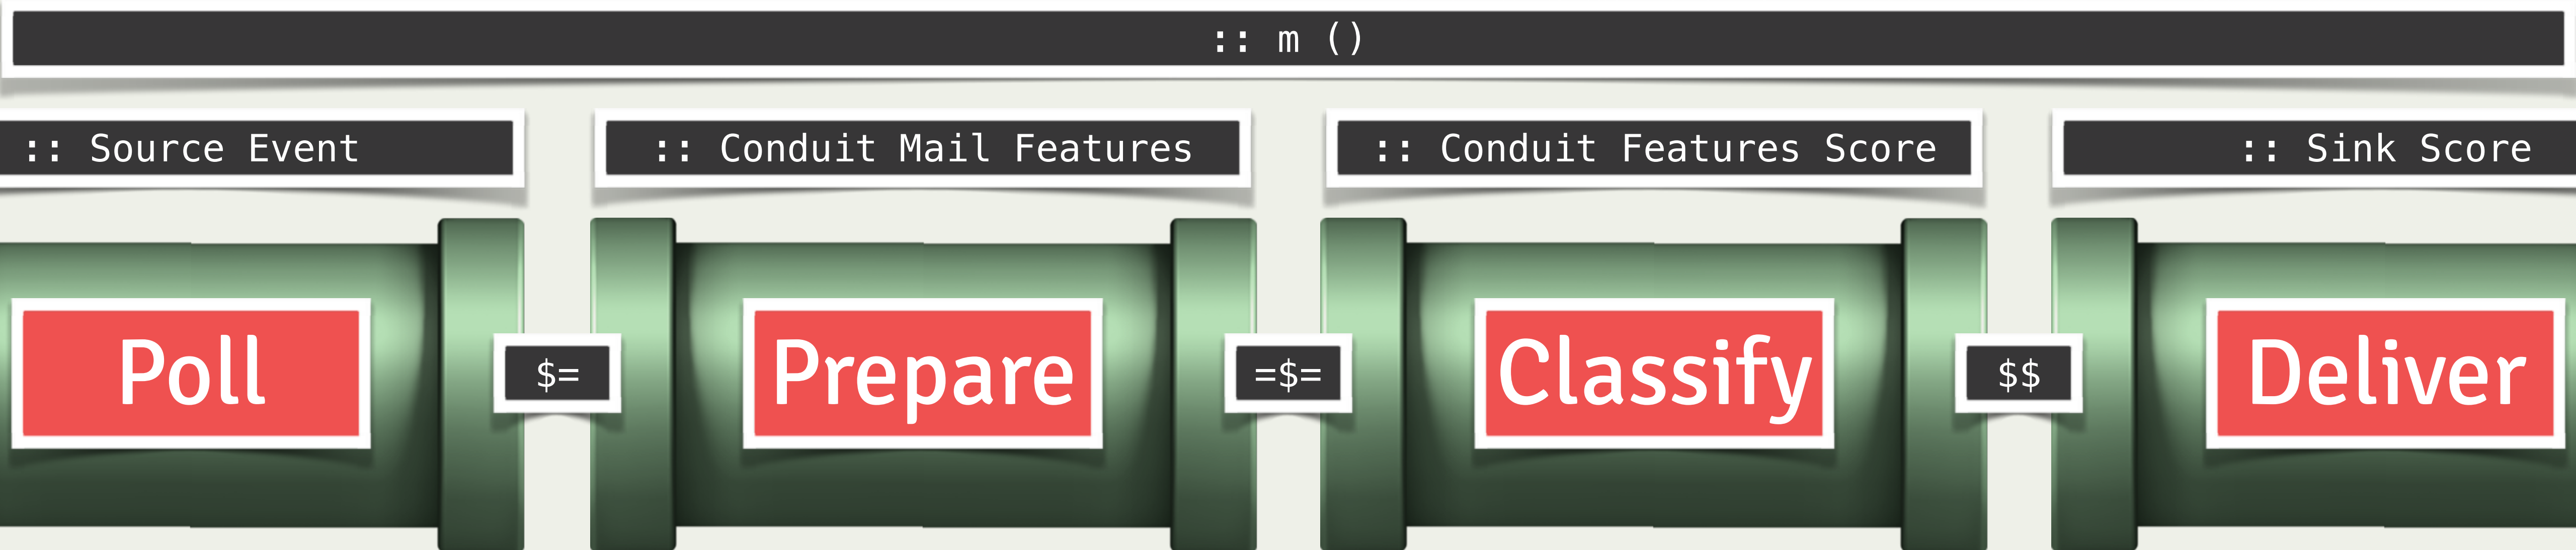




Conduit



Conduit'



Conduit''

```
:: Process m ()
```

```
: Process m Event
```

```
:: Process1 Mail Features
```

```
:: Process1 Features Score
```

```
:: Sink m Score
```

Poll

|>

Prepare

|>

Classify

to

Deliver

Scalaz Stream

```
run :: Process m a -> m ()
```

```
: Process m Event
```

```
:: Process1 Mail Features
```

```
:: Process1 Features Score
```

```
:: Sink m Score
```

Poll

|>

Prepare

|>




Classify

to

Deliver

Scalaz Stream'

Is Composition About  
Combinators or Laws?



# Vertical Composition



# Mail Delivery

Poll

Prepare

Classify


Deliver

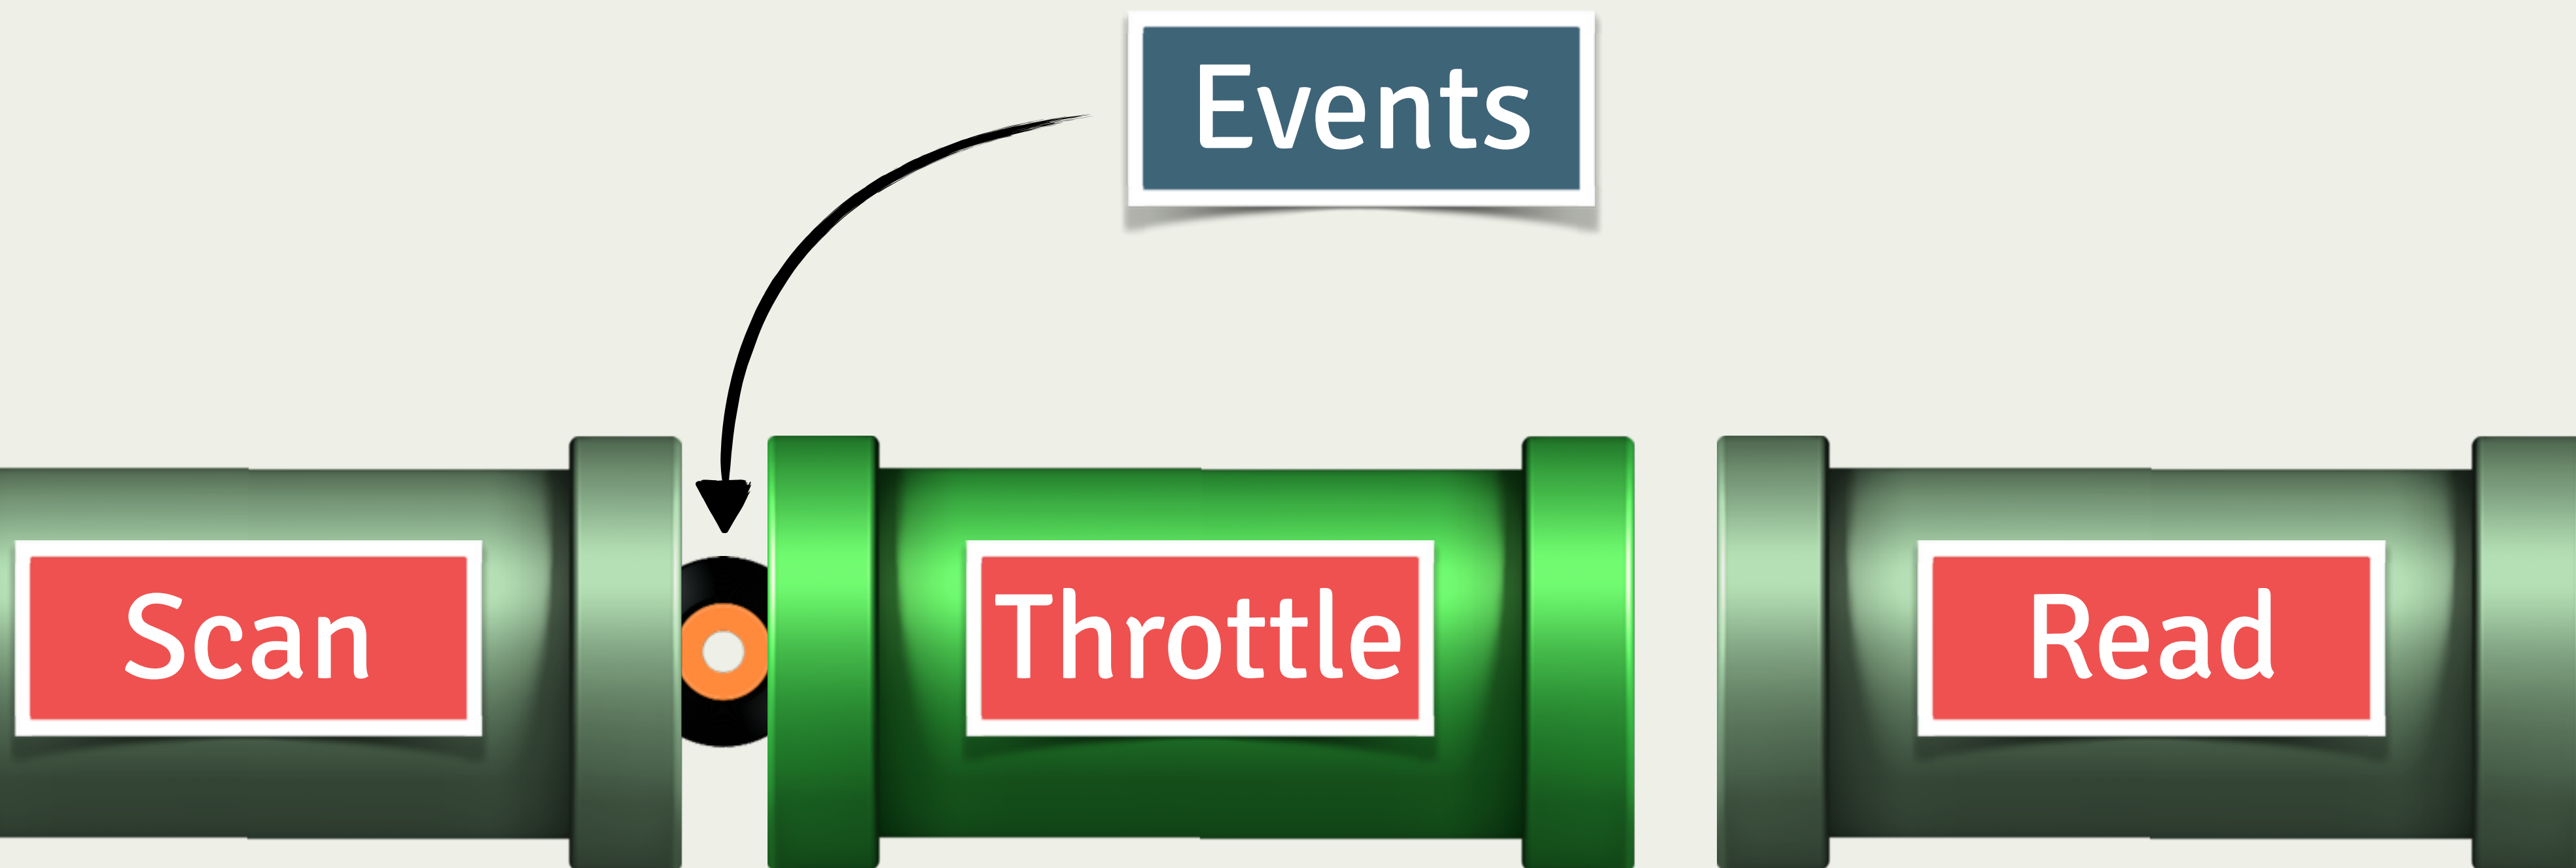
Poll

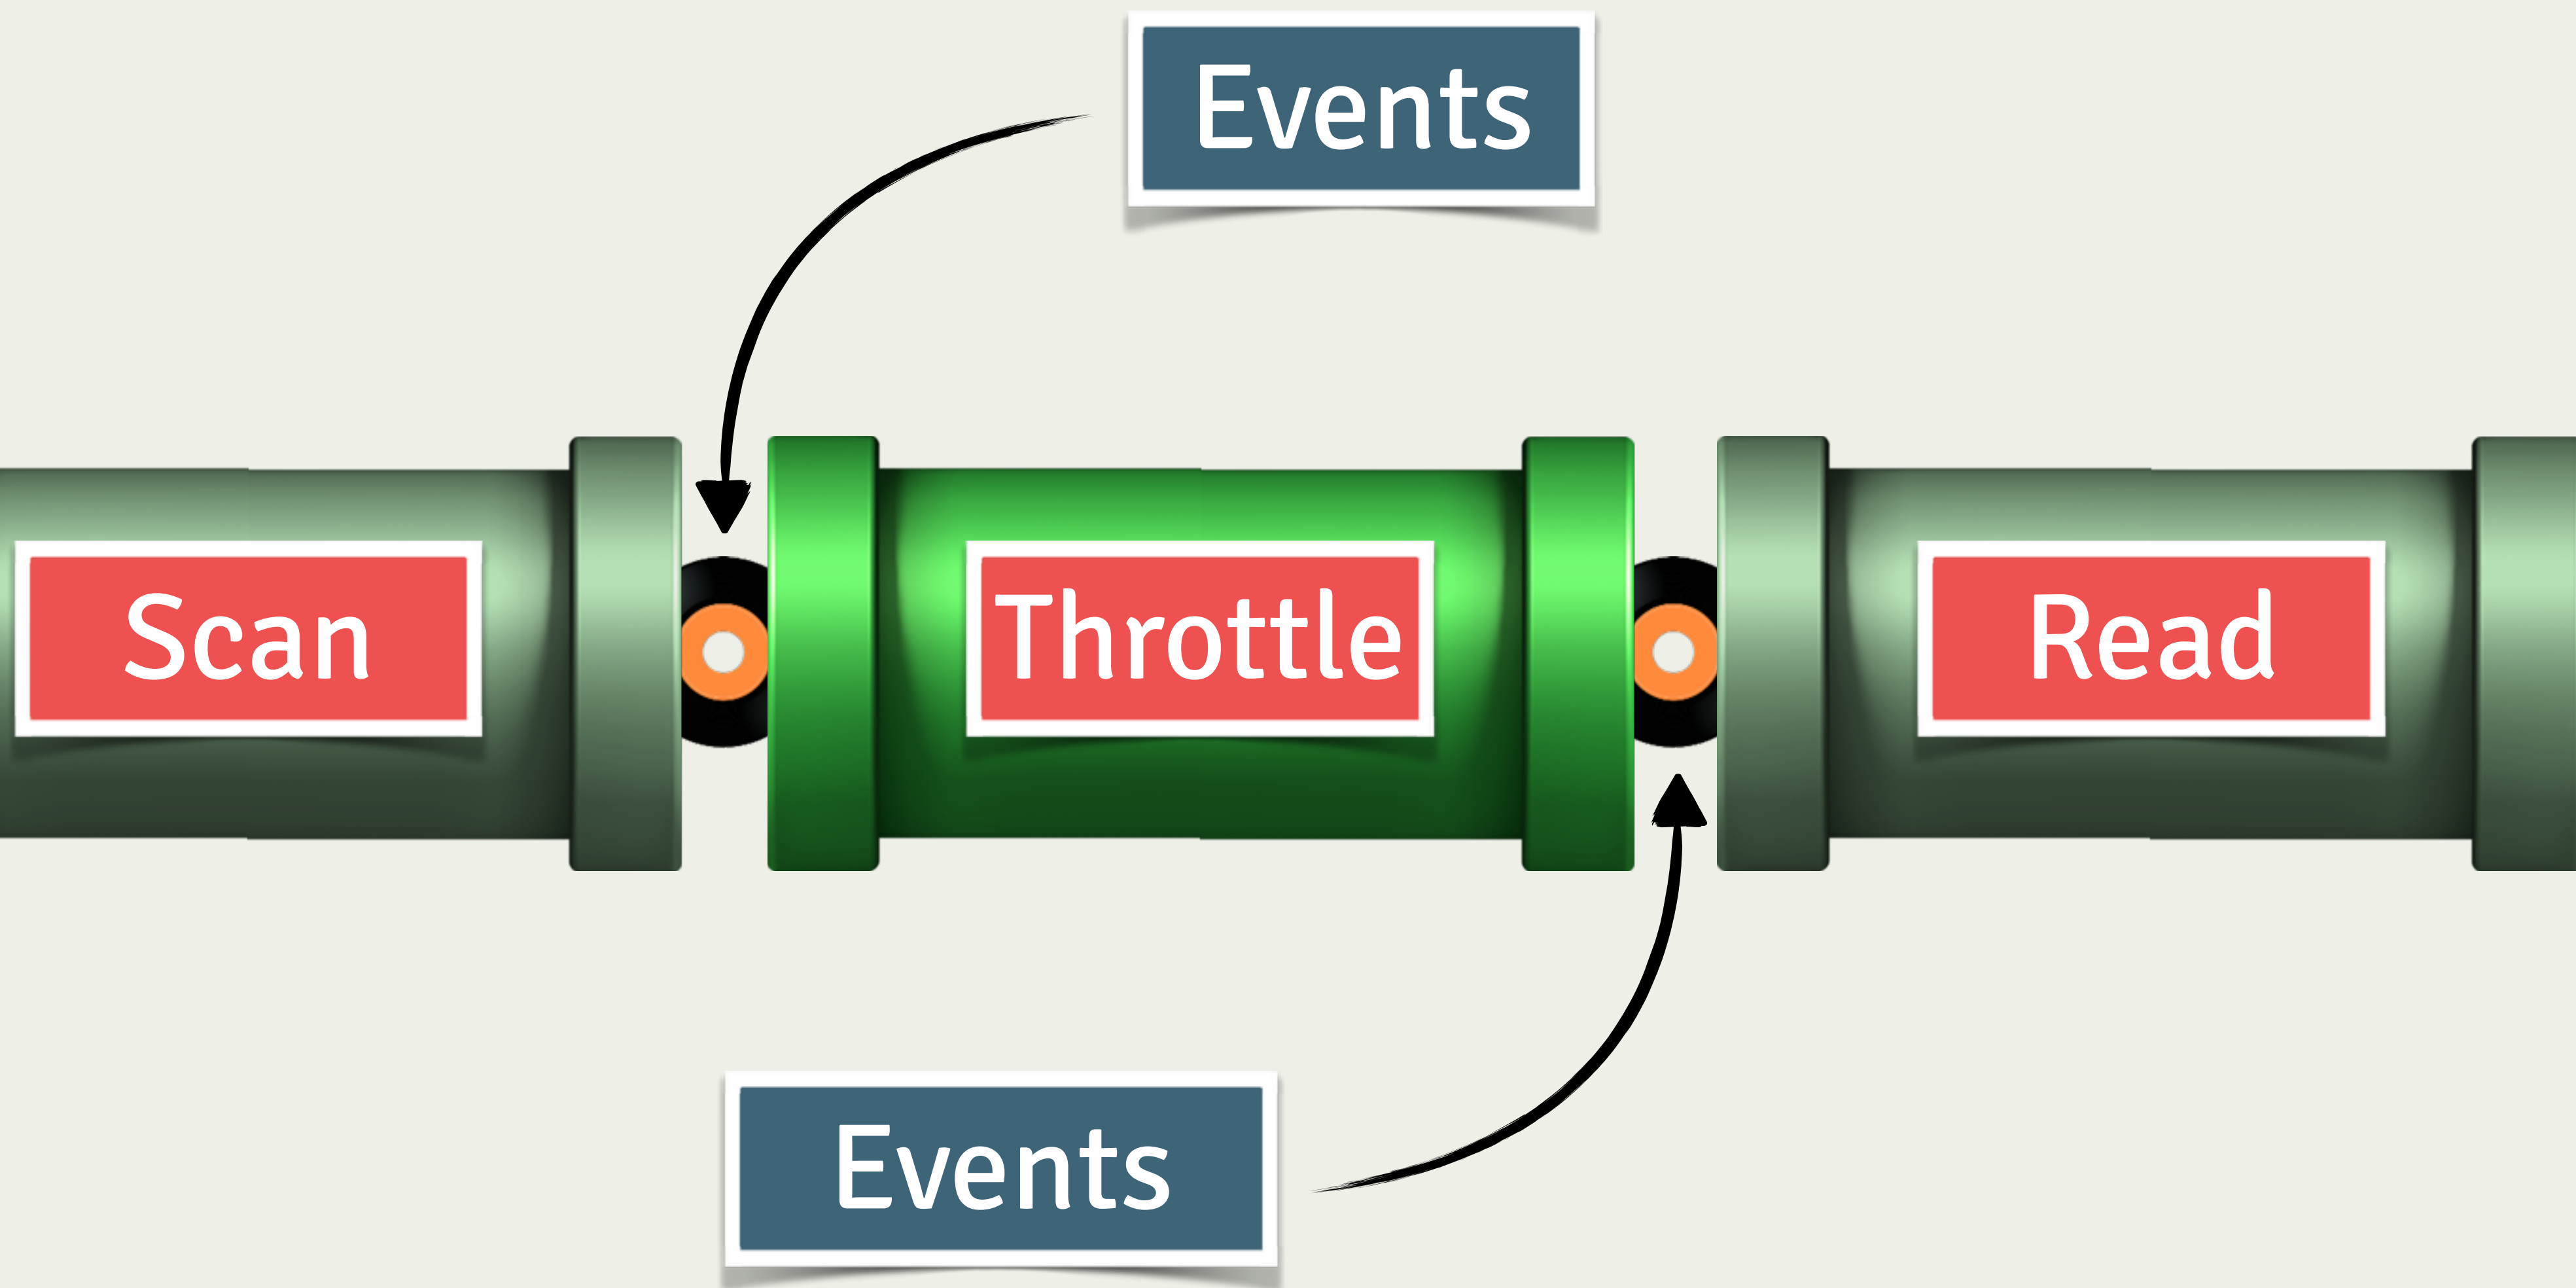
Scan

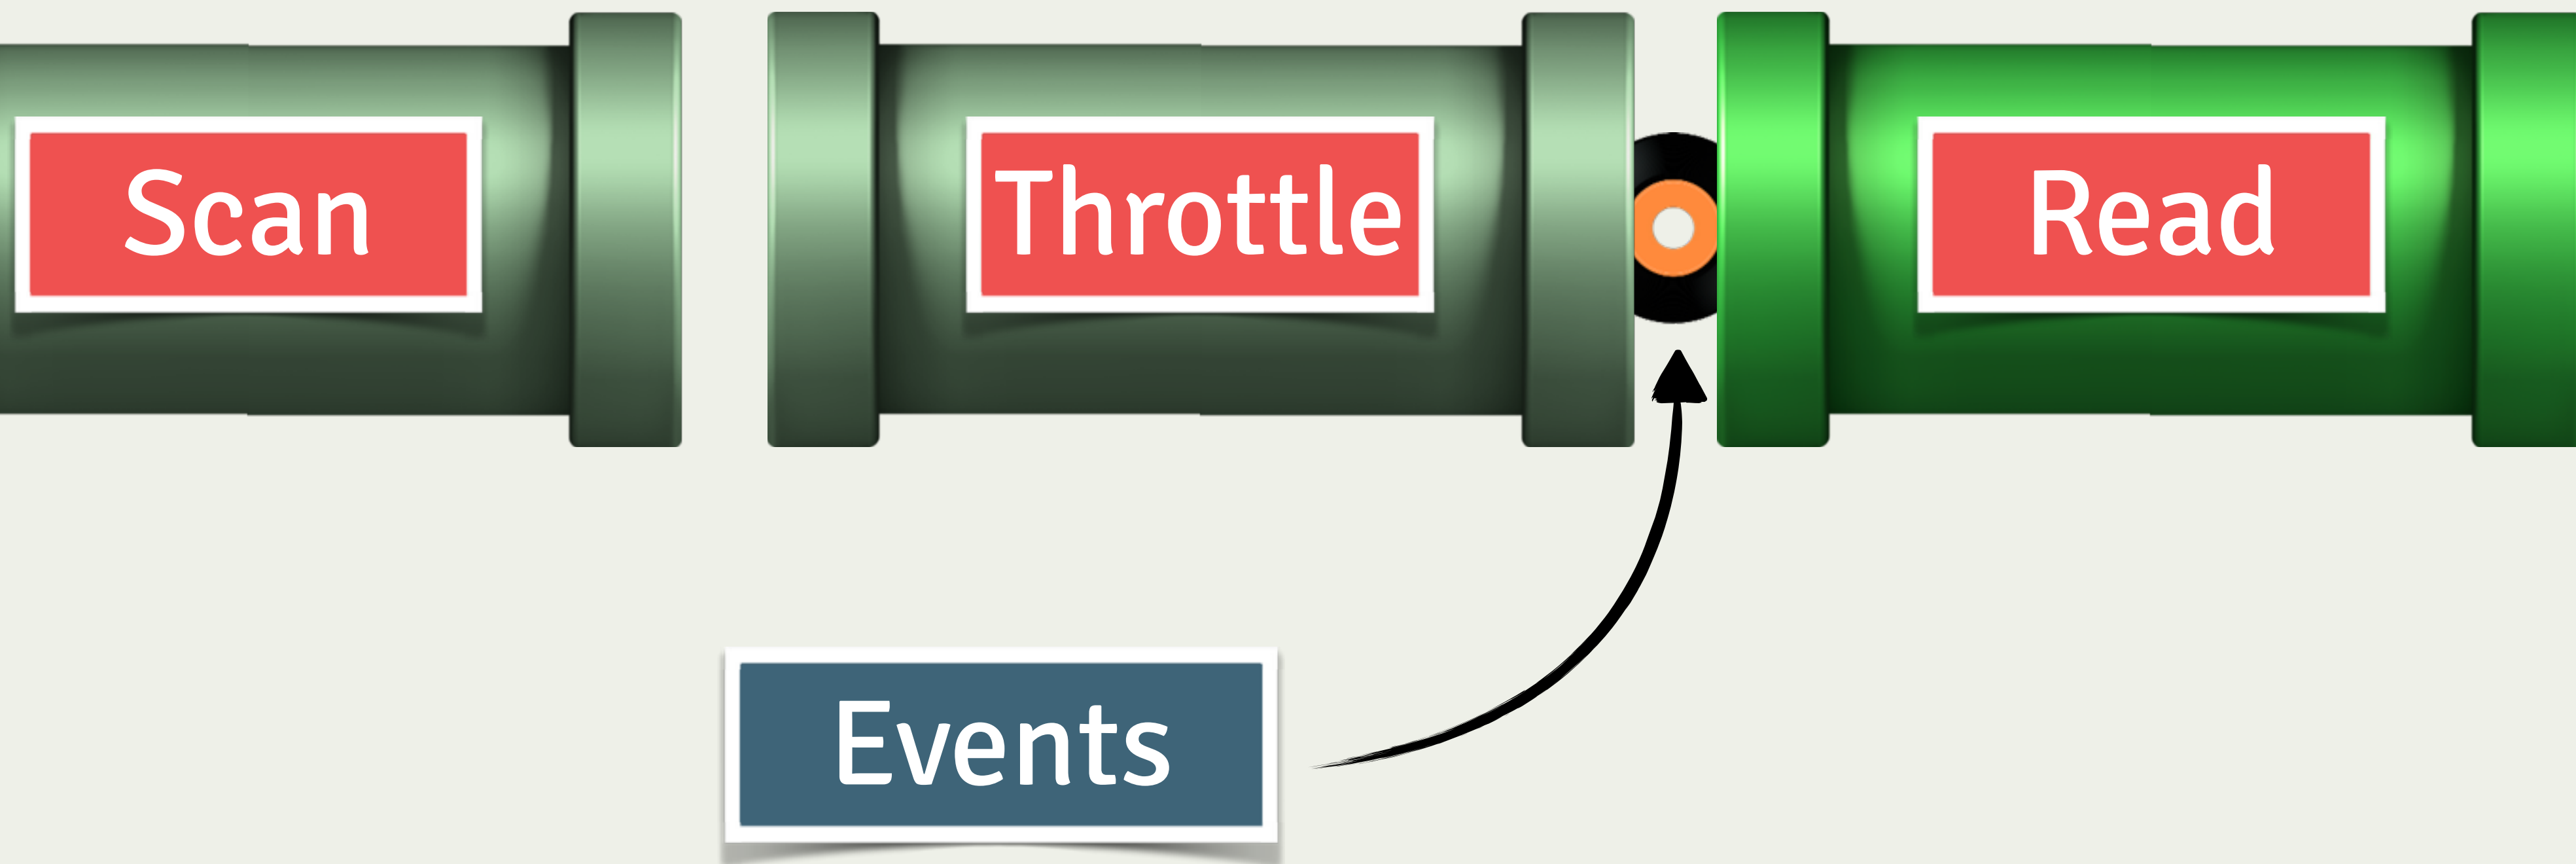
Throttle


Read





























































Pipes



Conduit



Scalaz Stream



# Intuition 3: Parsers

```
1 type In i m a = Pipeline i () m a
2
3 type Out o m a = Pipeline Void o m a
4
5 data Pipeline i o m a
6     = Done a
7     | Yield o (Pipeline i o a)
8     | Await (i -> Pipeline i o a)
9
```



```
1 type In i m a = Pipeline i () m a
2
3 type Out o m a = Pipeline Void o m a
4
5 data Pipeline i o m a
6   = Done a
7   | Yield o (Pipeline i o a)
8   | Await (i -> Pipeline i o a)
9
10 yield :: o -> Pipeline i o m ()
11 yield = Yield o (Done ())
12
13 await :: Pipeline i o m i
14 await = Await Done
```

```
1 one :: Pipeline i i m ()
2 one = do
3     i <- await
4     yield i
5
6 cat :: Pipeline i i m ()
7 cat = forever one
8
9 pairs :: Pipeline i (i, i) m ()
10 pairs = forever $ do
11     i1 <- await
12     i2 <- await
13     yield (i1, i2)
```

```
1 counter :: Monad m => Pipeline i (Int, i) m ()
2 counter = flip evalStateT 0 . forever $ do
3     i <- lift await
4     n <- get
5     lift . yield $ (n, i)
6
7 filter :: (i -> Bool) -> Pipeline i i m a
8 filter f = forever $ do
9     i <- await
10    when (f i) $ yield i
```

```
1 yield :: o -> Pipe i o m ()
2
3 await :: Pipe i o m i
```

## Pipes

```
1 yield :: o -> ConduitM i o m r
2
3 await :: ConduitM i o m (Maybe i)
4
5 awaitForever :: (\i -> ConduitM i o m a)
6               -> ConduitM i o m ()
```

## Conduit

```
1 emit :: o -> Process f o
2
3 await1 :: Process1 i i
```

## Scalaz Stream

**Subtlety** Fights Back

# Internal vs External Management of Resources

# Layered Streams

# Constant Memory Streaming



How much does elegance  
cost?

**to be continued...**

**@markhibberd**